I am an engineer. Most recently I've been working on evals for LLM agents; finding how (and why) deployed systems fail.

At BMO's AI Centre of Excellence I built evaluation harnesses and test setups that caught an agentic tool (serving over $200B in AUM) subtly downplaying investment risk. Behaviour that was invisible until I scaled up to hundreds of test cases to identify systematic misalignment. Now I'm designing evals for agents that retrieve and reason over commercial banking and insurance policy, and developing RL environments to train specialized agents for our wealth management division.

On the side I do research at Merlyn Labs, the 3-person research collective I co-founded. We built VLM judges that turn rollouts into dense, hard-to-game reward metrics, and placed 8th in the Stanford BEHAVIOR-1K challenge fine-tuning VLAs to accomplish household tasks in sim. I also worked on a flow-matching VLA integration for the RLinf framework so the BEHAVIOR-1K suite can be used for RL training in OmniGibson. Same category as Prime-RL or Verl.

I'm also training agents to play Settlers of Catan and Twilight Imperium. The interesting bit is taking a large action space and introducing multi-agent negotiation. Then seeing whether an agent can hold a long-term strategy across a whole game, and when it reaches for deception, blackmail, or betrayal to get there.

I'm a Canadian citizen and TN-eligible. Can relocate to NYC or SF within 4 weeks of an offer.

Shayan