

Shayan Shakeri

Toronto, Canada | shayansnezhad@gmail.com | +1 416 471 6651 | sshakerinezhad.github.io

By day, I work as an AI Research Engineer at BMO's AI Centre of Excellence. At night (and on weekends), I work at Merlyn Labs, the three-person research collective I co-founded. That mix is roughly what you're hiring for, and exactly how I like to work. Someone experienced in research who can apply it to all sorts of business cases.

At both, the through line of my work is designing evals and RL reward signals. Getting in the mud, finding where learned systems break, and fixing them. Things like:

- **Designing VLM judges that score rollouts into dense, context-dependent RL rewards**, built to be difficult to game, the core research thread at Merlyn Labs.
- **Uncovered systematic risk-downplaying bias in a production GenAI tool**, with a deterministic eval pipeline running hundreds of synthesized inputs against a \$200B+ AUM wealth-management system.
- **Doubled $\pi 0.5$'s position-swap success on LIBERO-PRO (21% to 42%)**, with a conservative finetuning recipe, after tracing the published checkpoint's brittleness to recipe-induced overfitting.
- **Open-sourced a flow-matching VLA integration for RLinf**, enabling RL training on the BEHAVIOR-1K suite in OmniGibson (8th place in Stanford's challenge, 10,000+ demonstrations across 22 tasks).
- **Built BardSong, an AI D&D storytelling tool**, fine-tuned an LLM on sourcebooks for narrative coherence; in closed alpha with 23 DMs.

Eval and reward signals are highly exciting for me: I see it as a sort of cat and mouse game between researcher and model. War gaming those interactions out and designing hard to game evals and RL environments.

My first move at Haize would be the one I run everywhere: probe where the model or the metric breaks, make the reward signal expensive to game, and measure against a baseline that's actually calibrated. I haven't contributed to Verl or Prime-RL yet, but RLinf is the same class of infra, so the ramp is short.

I'm a Canadian citizen, TN-eligible with no sponsorship needed, and happy to be in New York or San Francisco within 2-4 weeks of an offer.